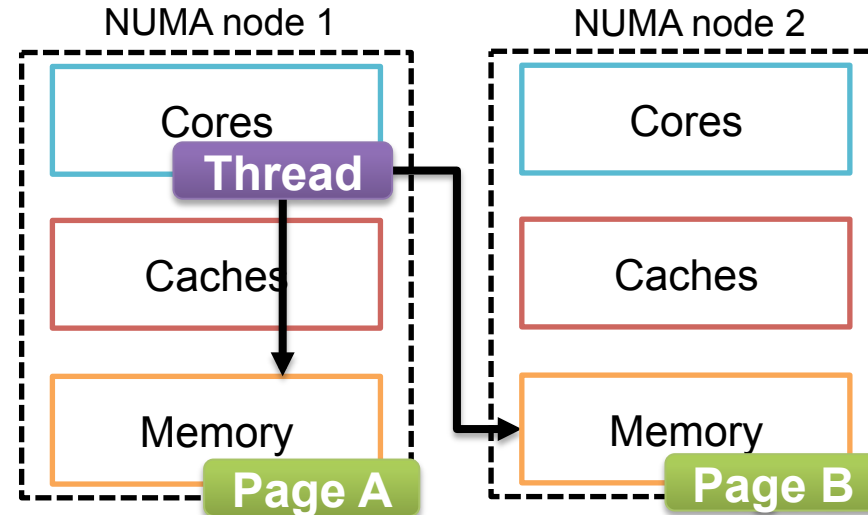


# Locality vs. Balance: Exploring Data Mapping Policies on NUMA Systems

Matthias Diener, Eduardo Cruz, Philippe Navaux

*Federal University of Rio Grande do Sul (UFRGS), Porto Alegre, Brazil*

- Shared-memory architectures with **multiple memory controllers (MC)**
  - Non-uniform memory access (**NUMA**) behavior
  - Each controller forms a NUMA node, can access its own part of the main memory



**Data mapping:** map memory pages to NUMA nodes to improve memory access performance

- Basic data mapping policy: **first-touch**
  - Map page to **first** node that accesses it
- Types of improved policies
  1. **Locality-based** mapping
    - Traditional policy to improve data mapping
    - Minimize remote memory accesses
    - Map pages to the node with the **most** accesses
  2. **Balance-based** mapping
    - Recent policies for modern NUMA systems
    - Avoid overloading controllers
    - Map pages in such a way that the load is balanced

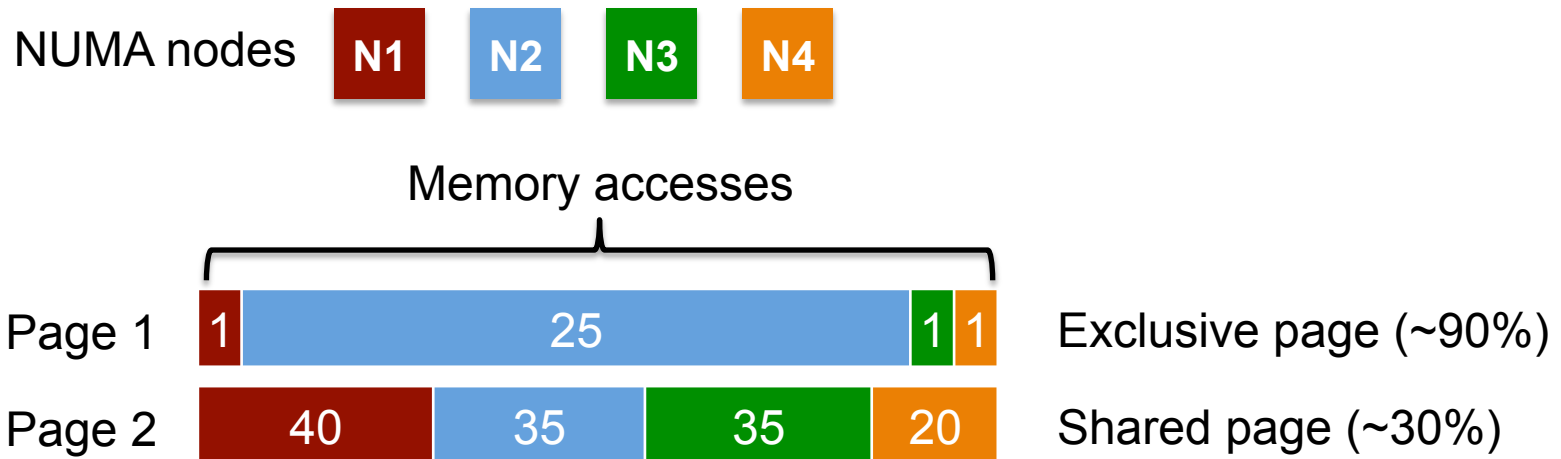
1. Which memory access behavior can be improved by mapping?
2. Which improvements can be achieved with different policies?

1. **Metrics** to describe memory access behavior of parallel applications.
2. **Mapping policies** that improve locality, balance or both.
3. **Evaluation** of a large set of parallel applications on several NUMA platforms.

# **Describing the memory access behavior of parallel applications**

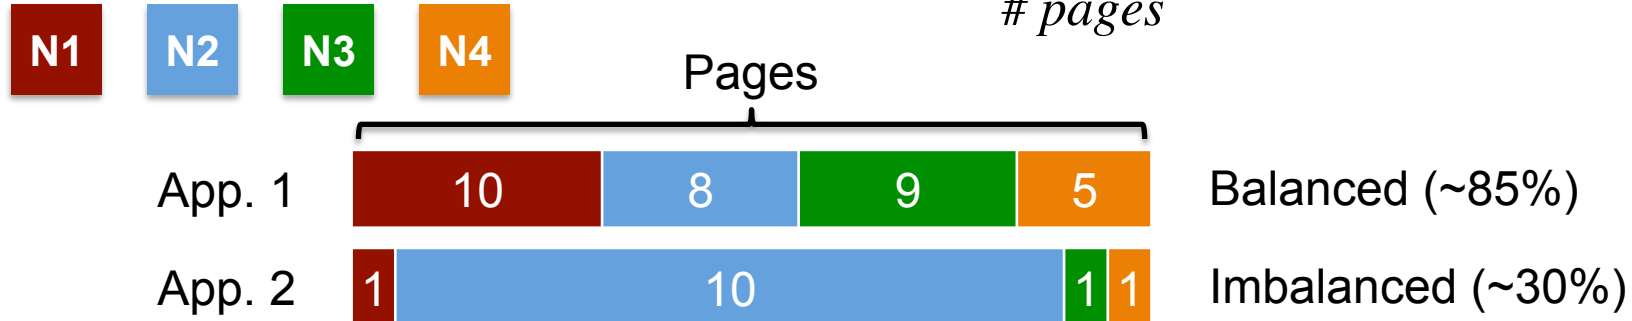
- Two sets of metrics:
  - **Exclusivity** (locality-based mapping)
    - Property of application
  - **Balance** (balance-based mapping)
    - Property of application and data mapping policy
- Describe per-node memory access behavior
- Granularity: memory page size
  - Currently: several KByte (4 KByte on x86 by default, our baseline)
  - Future systems may have larger pages

- **Page Exclusivity:** highest % of memory accesses from same NUMA node




- **Application Exclusivity:** Calculate weighted exclusivity for all pages:  
Page 1: 28 accesses, Page 2: 130 accesses  
 $(28 \cdot 0.9 + 130 \cdot 0.3) / 2 = 64\%$
- Min:  $1/\text{\#nodes}$  (fully shared), Max: 100% (fully exclusive)
- Expect higher improvements from locality-based mapping with high exclusivity

- Page balance**  $1 - \frac{\max(\# \text{ pages\_node}) - \min(\# \text{ pages\_node})}{\# \text{ pages}}$



- Memory access balance:** calculate weighted balance



- Min: 0% (all pages/accesses to single node), Max: 100% (fully balanced)
- Expect more improvements from balanced-based policy with low balance



# Mapping policies

- Policies
  - **First-touch**
  - **Random, Interleave** (independent of memory access behavior)
  - **Locality, Balanced, Mixed** (depend on memory access behavior)
- First-touch
  - Page is mapped to NUMA node with first memory access to page
  - Default policy of most OS (baseline)
- Random
  - Assign pages to nodes randomly
  - Validate importance of mapping
- Interleave
  - Traditional policy to balance pages between nodes
  - $\text{node}[p] = \text{address}(p) \bmod \text{\#nodes}$

- Locality
  - Traditional policy to maximize memory access locality
  - $\text{node}[p] = \text{node\_with\_most\_accesses}(p)$
- Balanced
  - Balance number of memory accesses per NUMA node
  - Algorithm
    - Sort pages by number of memory accesses
    - Place each page on the node with the most accesses to it, that is not overloaded yet
  - Complementary policy to Locality
    - Favor balance over locality vs. favor locality over balance

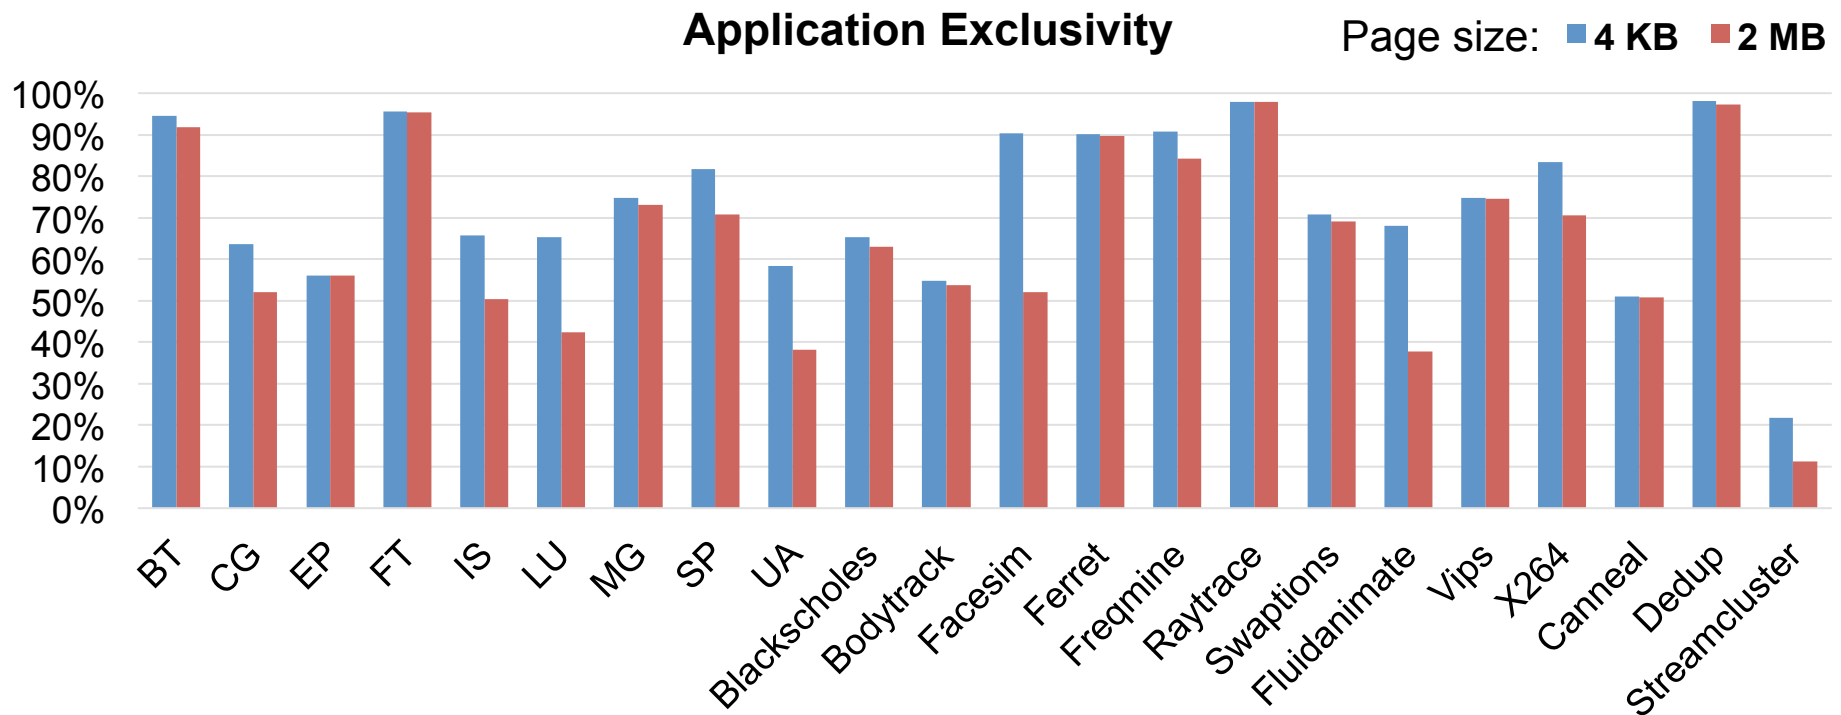
- Mixed
  - Apply locality policy for highly exclusive pages
  - Shared pages: distribute with Interleave policy

$$\text{node}[p] = \begin{cases} \text{node\_with\_most\_accesses}(p), & \text{if exclusivity} > \text{threshold} \\ \text{node}[p] = \text{address}(p) \bmod \# \text{nodes}, & \text{otherwise} \end{cases}$$

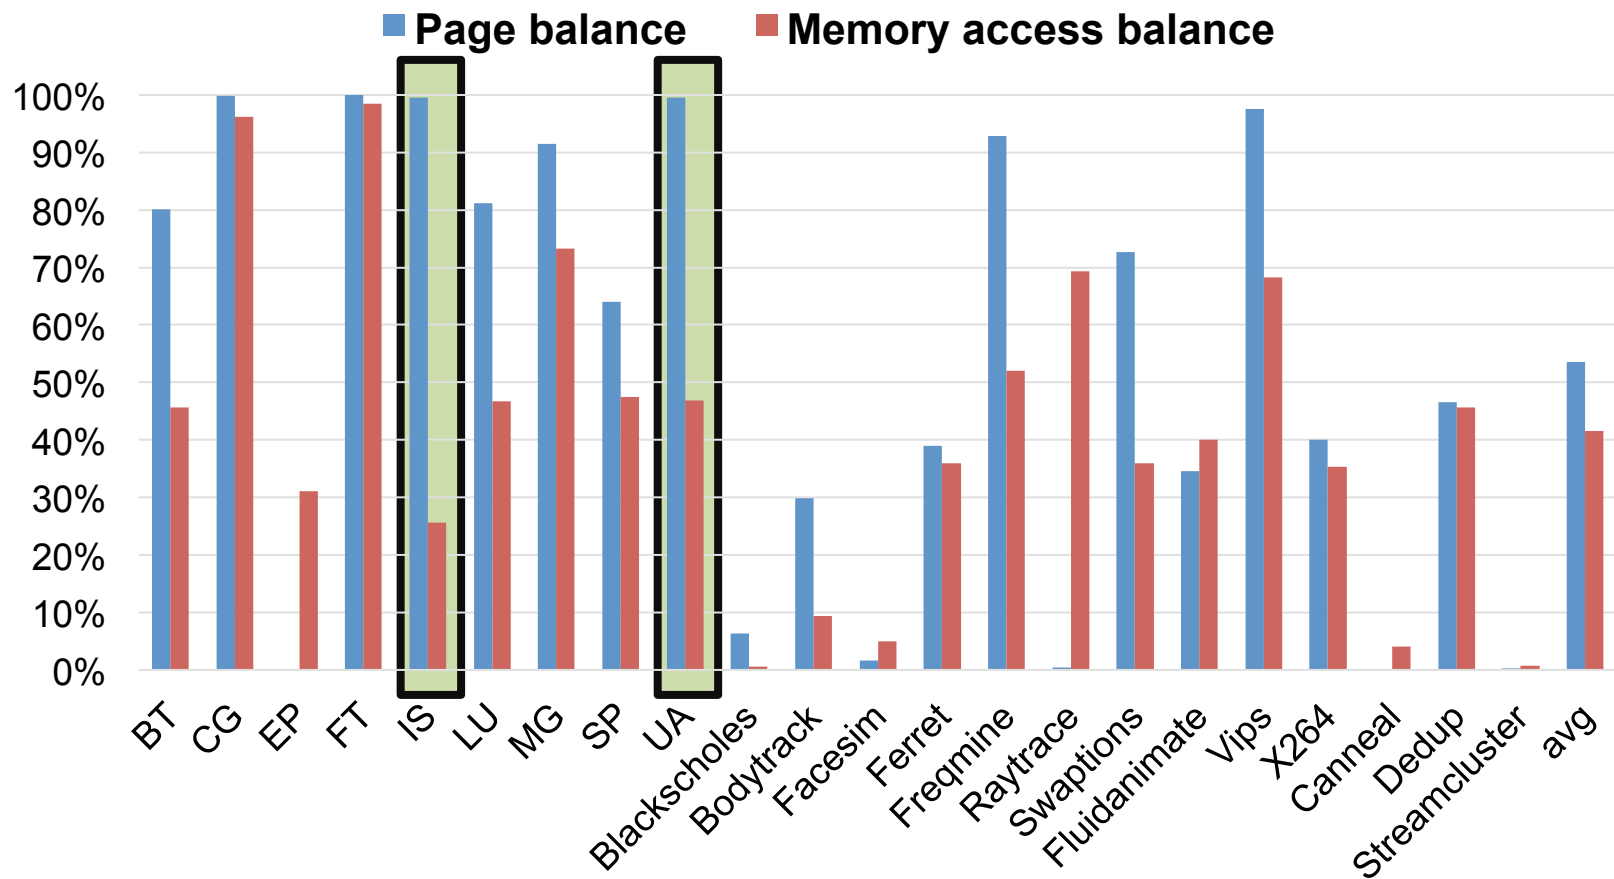
- Trade-off between Locality and Balance
- Evaluated various thresholds, 90% gave highest improvements

# Benchmark metrics

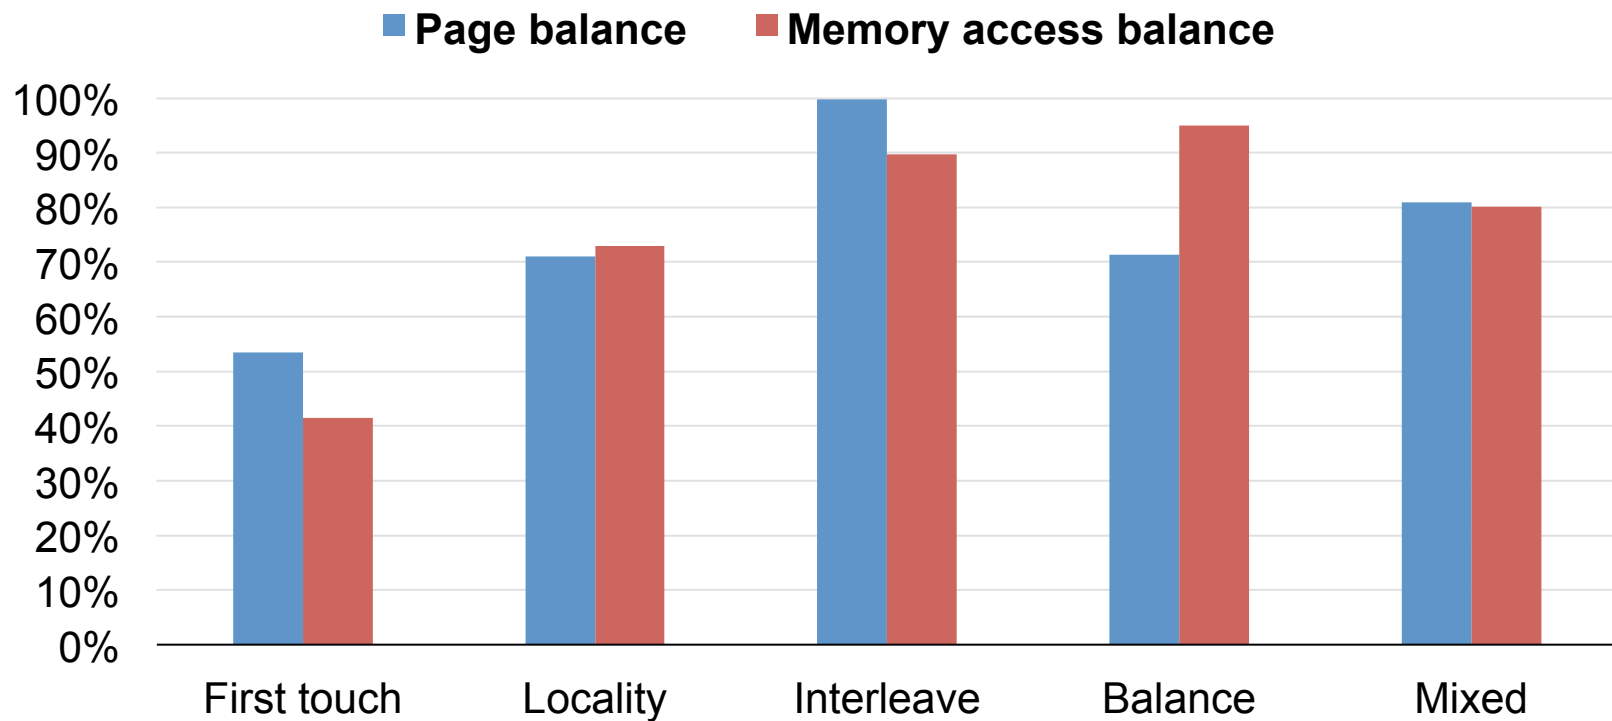
- Evaluate memory access behavior with memory tracing tool built with Pin
- Benchmark suites
  - NAS Parallel Benchmarks, OpenMP version (C input size)
  - PARSEC (*native* input size)
- 64 threads; 4 NUMA nodes for balance metrics
- Behavior is constant with same input data/number of threads
- Measure
  - Application exclusivity for 2 page sizes
  - Page and memory access balance for all mapping policies



Average: 73% (4KB), 65% (2MB)







- Exclusivity high for most applications
  - Suitable for locality-based data mapping
  - Decreases only slightly for larger pages
- Most applications imbalanced
  - First-touch from master thread
  - Suitable for balance-based data mapping
  - Locality already balances better than first-touch


# Performance Evaluation

- 3 NUMA machines: Itanium, Xeon, Opteron
  - NUMA factor:  $\text{latency}(\text{remote access}) / \text{latency}(\text{local access})$


Machine	# NUMA nodes	# PUs	Processors	Page size	NUMA factor
Itanium	2	8	4x Itanium 9030, 2 cores	16 KB	2.1
Xeon	4	64	4x Xeon X7550, 8 cores, SMT	4 KB	1.5
Opteron	8	64	4x Opteron 6386, 8 cores, SMT	4 KB	2.8

- Same benchmarks/input sizes as before
  - Use memory tracer + offline data mapping algorithm
  - No benchmarks whose memory addresses change between executions
- Execute with 8/64/64 threads
- Compare all mapping policies
  - Performance improvements to first-touch mapping


Itanium: 2 NUMA nodes, high NUMA factor



Xeon: 4 NUMA nodes, low NUMA factor



Opteron: 8 NUMA nodes, high NUMA factor



- First-touch has a negative impact on performance
  - Even random page assignment is better
- Locality is still more important than balance
  - Results of Locality > Balanced
- Mixed policies can achieve the highest improvements
- What can be done?
  - Application developers
    - Prepare application for first-touch
    - Access pages only from single thread to increase exclusivity
  - OS developers
    - Use mechanisms to refine data mapping during execution
      - Basics already done (NUMA balance for Linux, ...)



# Conclusions

- Data mapping has a substantial influence on the performance of NUMA machines
  - Gains expected to rise in the future
  - Need to choose correct policy
- Memory access locality is most important metric to improve
  - Even on current NUMA architectures
- Policies that combine Locality and Balance can result in the highest improvements
- Future work
  - Apply policies to online mechanisms (no need for profiling step before execution)

# Locality vs. Balance: Exploring Data Mapping Policies on NUMA Systems

Matthias Diener, Eduardo Cruz, Philippe Navaux

*Federal University of Rio Grande do Sul (UFRGS), Porto Alegre, Brazil*

`mdiener@inf.ufrgs.br`



