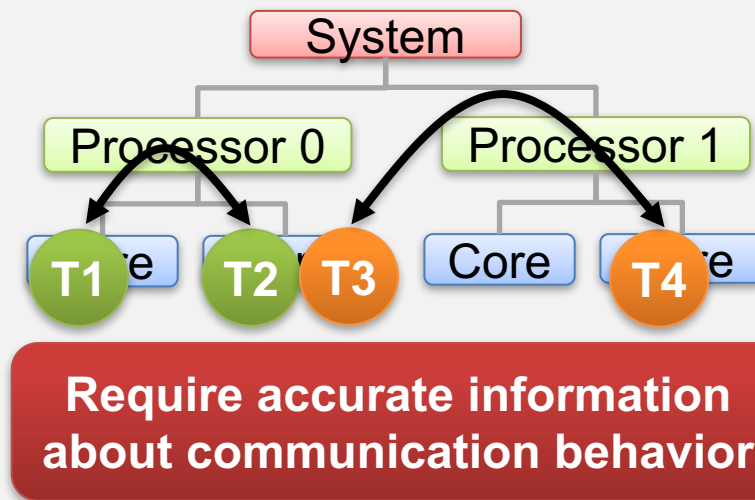


# Communication in Shared Memory: Concepts, Definitions, and Efficient Detection

Matthias Diener, Eduardo Cruz, Marco Alves, Philippe Navaux

*Federal University of Rio Grande do Sul, Porto Alegre, Brazil*

- Parallel applications need to communicate
- Communication has a high overhead
  - Lower performance and energy efficiency compared to computation
- **Solution: Map tasks such that communication impact is reduced**




- In shared memory
  - Communication through memory accesses
  - Mapping improves interconnection and cache usage

## Distributed memory

**Explicit** communication (MPI, Charm++, ...)

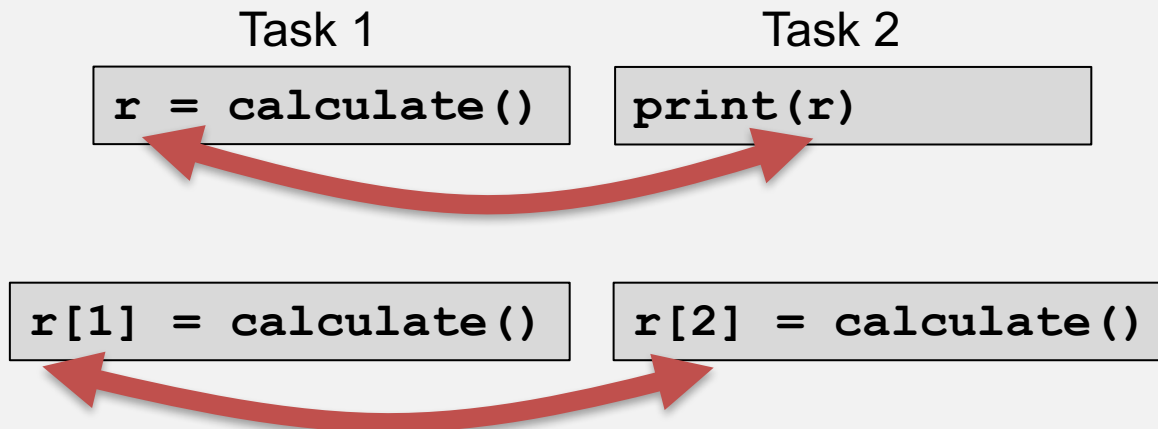
- Function calls for communication
- Detection via function instrumentation



## Shared memory

**Implicit** communication (OpenMP, Pthreads, ...)

- Communication via memory accesses to shared data
- Detection via memory access traces



- Detection of communication in shared memory not straightforward
  - Nevertheless, need accurate information to perform mapping
  - Focus: architectural impact of communication to perform the mapping
    - Sensitive to architectural parameters
      - Cache size, line size, ...

## How to describe implicit communication on the architectural level?


1. Types of communication
2. Detecting communication accurately and efficiently
3. Impact on performance gains

# 1. Communication types


- In **explicit** communication, all communication is **true**
  - Intention to exchange data
- Not the case in **implicit** communication
  - Not every memory access to shared data is intention to communicate
  - "Unintentional" (**false**) communication
    - **Spatial, logical, temporal**
  - Caused by operation of cache subsystem

Example:






**True**




**Spatial false (false sharing)**

**Logical false**



**Temporal false**



- Impact on mapping
  - **True** comm.
  - **Spatial false** comm.
  - **Logical false** comm.

} Consider for mapping
- No impact
  - **Temporal false** comm.

} Do not consider for mapping


## Communication event:

Two memory accesses from different threads to the same cache line while the cache line is not evicted.

**Accurate** definition of communication




- **Direction** of communication
  - **Explicit** communication: well-defined sender/receiver
  - No direction in **implicit** communication
    - E.g., two tasks read the same data
- Create **communication matrix** with communication events



**Communication matrix determines the mapping.**


- Benchmarks
  - NAS Parallel Benchmarks (NPB), OpenMP implementation; *A* input
  - PARSEC (Pthreads, OpenMP); *simlarge* input
  - Select 5 benchmarks with different behaviors
    - LU, UA
    - Blackscholes, Ferret, Swaptions
- Microarchitecture simulator
  - Sandy Bridge: 4x 8-core processors (2-SMT), 64 threads
    - 4x 18 Mbyte L3 cache (inclusive)
- Measure
  - **True, spatial false, temporal false** communication
    - No logical false communication, included in other types
    - Temporal false communication: simulate infinite cache




- Little spatial false communication
- High temporal false communication (esp. LU/Ferret)
  - Does it change the pattern?

**Accurate**



**LU**



**Ferret**



**Accurate +  
temp. false**






## **2. More efficient detection**

- Disadvantages of previous mechanism
  - Full cache simulator
    - High runtime overhead
  - Cache line granularity
    - High storage overhead
- Provide more efficient ("**relaxed**") definition
- Main idea: cache hierarchy mostly not relevant
  - Want to detect communication between **tasks**
  - Necessary to remove temporal false communication
    - Need simpler way to filter
- Approximate temporal information via short queues

- Divide address space into memory blocks (block size  $\geq$  line size)
- Maintain small FIFO queue (2 positions) of tasks that accessed the block
  - Detect communication within the queue
- For each memory access (**addr**, **tid**):
  1. Find memory block: **addr**  $\gg$  **granularity**
  2. Number of other tasks that previously accessed block ( $\leq 2$ )

| No access                                 | 1 access  | 2 accesses  |
|---|---|---|
| Enqueue <b>tid</b> .<br>No communication. | Enqueue <b>tid</b> .<br>Communication between<br>old and new <b>tid</b> . | Enqueue <b>tid</b> , dequeue<br>oldest.<br>Communication between<br>new <b>tid</b> and both old ones. |




- Overhead
  - Constant time operation per memory access
  - No expensive operation (cache simulation, time stamps, ...)
- Granularity/block size (default: 64 Bytes)
  - Larger: more spatial false communication, less storage required



- UA benchmark (highest error), 8 threads


**Accurate  
definition**




**Relaxed  
definition**

Granularity


**64 Byte**




**1 KByte**



**16 KByte**




**64 MByte**



### **3. Performance gains**

- Methodology
  - Same benchmarks as before, 64 threads
  - Real machine: Intel Sandy Bridge, 4x 8-core processors (2-SMT)
    - L1/L2 caches per core, L3 cache per processor
  - Compare to default Linux 3.8 scheduler (CFS)
  - Thread mappings calculated with Scotch library
    - Static, no runtime overhead
- Compare execution time with mapping
  - Accurate definition
  - Accurate + temporal false communication
  - Relaxed definition with different granularities (64 Byte, 1Kbyte, 1Mbyte)
- Compare overhead accurate vs. relaxed communication detection



## Overhead

Slowdown vs. normal execution

| Mechanism | LU     | UA     | Blackscholes | Ferret | Swaptions |
|-----------|--------|--------|--------------|--------|-----------|
| Accurate  | 5944 x | 2860 x | 1771 x       | 5304 x | 6157 x    |
| Relaxed   | 49 x   | 72 x   | 113 x        | 39 x   | 148 x     |

### Available at

CacheSim: <http://github.com/matthiasdiener/CacheSim>

Numalize: <http://github.com/matthiasdiener/numalize>

Both are based on Intel Pin.

# Conclusions

- Accurate view of communication remains important challenge for thread mapping in shared memory
  - Can have huge impact on gains
  - Main challenge: temporal false communication
- Solutions: cache simulator, queue-based approximation
- Approximation solution
  - Similar gains as cache simulator (22% vs. 19%)
  - Substantially lower overhead (up to 100 times faster)
- Future work: impact of sampling

**Thank you!**

# **Communication in Shared Memory: Concepts, Definitions, and Efficient Detection**

Matthias Diener, Eduardo Cruz, Marco Alves, Philippe Navaux

*Federal University of Rio Grande do Sul, Porto Alegre, Brazil*

**[mdienner@inf.ufrgs.br](mailto:mdienner@inf.ufrgs.br)**

CacheSim: <http://github.com/matthiasdiener/CacheSim>  
Numalize: <http://github.com/matthiasdiener/numalize>